

Representation-Steered Incremental Adapter-Tuning for Class-Incremental Learning with Pre-Trained Models

Jiarui Zhao^{1,2}

Libo Huang^{1*}

Xiangqi Li^{1,3}

Zhulin An^{1*}

Chuanguang Yang¹

Yu Wang^{1,3}

Boyu Diao¹

Yongjun Xu¹

¹State Key Lab. of AI Safety, Institute of Computing Technology, Chinese Academy of Sciences

²School of Advanced Interdisciplinary Sciences, University of Chinese Academy of Sciences

³School of Computer Science and Technology, University of Chinese Academy of Sciences

zhaojiarui231@mails.ucas.ac.cn {huanglibo, anzhulin}@ict.ac.cn

{lixiangqi24s, yangchuanguang, wangyu23s, diaoboyu2012, xyj}@ict.ac.cn

Abstract

Class-Incremental Learning (CIL) aims to develop models to continuously learn new classes without forgetting learned old ones. Recent advances combine pre-trained models with parameter-efficient fine-tuning, achieving promising results. However, these approaches typically allocate new trainable parameters for each task, causing the model size to grow linearly with task number. Moreover, they lack explicit mechanisms to structure a coherent and discriminative representation space across tasks. To address these limitations, we propose Representation-Steered Incremental Adapter Tuning (RSIAT). RSIAT maintains a single shared adapter for all tasks, eliminating parameter growth during incremental learning. In the base task, we introduce a representation-steering loss that enhances discriminative feature learning while facilitating future task adaptation. During incremental tasks, a residual autoencoder-based projector aligns feature distributions between old and new tasks, preserving representation consistency without over-constraining the shared adapter. Extensive experiments on six CIL benchmarks demonstrate that RSIAT significantly outperforms state-of-the-art methods in both performance and parameter efficiency, achieving superior stability-plasticity trade-offs with minimal trainable parameters. Code is available at <https://github.com/zjrjrz/RSIAT>.

1. Introduction

In traditional deep learning, models can access all data for unified training. However, real-world data often arrives in streaming formats. Therefore, real-world vision systems are required to continuously learn new classes without for-




Figure 1. Parameter-Accuracy comparison on ImageNet-R (B0I20). Each point shows the average accuracy (%) versus the number of total trainable parameters (Million). All methods are based on the same pre-trained model (ViT-B/16-IN21K). Our method (RSIAT, red star) uses a constant-size adapter, avoiding task-specific module growth while further improving accuracy.

getting old ones [2], known as Class-Incremental Learning (CIL). In CIL, a major challenge is addressing catastrophic forgetting [8], and various methods [18, 23, 24] have been developed to alleviate this issue.

Traditional CIL methods can be broadly categorized into replay-based [3, 4, 14, 32, 64], regularization-based [1, 15, 34, 58], and isolation-based methods [26, 27, 33]. However, with the widespread adoption of large Pre-Trained Models (PTMs) [11], research has gradually shifted from “training from scratch” to Parameter-Efficient Fine-Tuning (PEFT) [16, 21, 22, 42–44]. Unlike updating the entire backbone network, PEFT adapts small auxiliary training modules (e.g., prompts [17], adapters [5]) to the frozen

*Corresponding authors.

backbone. PTM-based CIL methods have greatly advanced performance but also exposed two fundamental issues. First, most approaches incrementally accumulate task-specific modules as new tasks arrive. When a new task is introduced, the model typically appends additional prompts or adapters that must be retrieved during inference. This design causes a linear increase in trainable parameters, and mis-selection or interference among modules can significantly degrade performance [10]. Second, despite the strong generalization capability of PTMs, current methods offer limited guidance for constructing a well-structured and task-consistent representation space. Without explicit regularization [49, 50], newly learned tasks induce feature drift in previously learned categories, manifesting as prototype shifts [7, 45], blurred decision boundaries [56], and reduced representation stability [62].

Motivated by these issues, we posit that an ideal CIL system should satisfy two key requirements: (i) parameter growth should be limited or ideally constant throughout training; and (ii) the feature space should remain compact within classes, sufficiently separable across classes, and aligned between historical and new tasks. To this end, we propose Representation-Steered Incremental Adapter Tuning (RSIAT), a PTM-based CIL framework designed from a representation-first perspective. RSIAT shares a single adapter across all tasks, preventing parameter bloat and eliminating the need for instance-level module retrieval. In the base task, we introduce a representation-steering loss that encourages intra-class compactness and inter-class separability while preserving PTM transferability. This loss promotes discriminative structures, such as well-defined decision boundaries and prototype compactness, enabling subsequent tasks to build on a robust and adaptive feature space. In incremental tasks, a lightweight residual autoencoder-based projection maps new task features to the representation space established by previous tasks. This residual alignment mechanism suppresses prototype drift, maintains clear decision boundaries, and improves representation stability without over-regularizing the shared adapter. As illustrated in Fig. 1, RSIAT achieves higher average accuracy with fewer trainable parameters than existing PTM-based CIL methods, striking a favorable balance between accuracy and parameter efficiency.

In summary, our main contributions are:

- We propose RSIAT, a PTM-based CIL framework that shares a single adapter across all tasks, preventing parameter growth while enabling efficient incremental learning.
- We introduce a representation-steering loss in the base task and a residual autoencoder-based projection in incremental tasks to enforce compact, discriminative, and aligned feature spaces, mitigating prototype drift and preserving decision boundaries.
- We demonstrate that RSIAT achieves state-of-the-art

accuracy-efficiency trade-offs on six CIL benchmarks, outperforming existing methods while keeping trainable parameters constant.

2. Related Work

2.1. Class-Incremental Learning (CIL)

Class-Incremental Learning (CIL) extends continual learning by requiring models to continuously learn new classes while maintaining performance across all previously encountered classes under a unified classifier [61]. Existing CIL approaches can be broadly categorized into three paradigms. Replay-based methods [3, 4, 32] mitigate forgetting by retaining exemplars or synthetic samples from past tasks and interleaving them with new data during training. These methods often design sophisticated sample selection or memory management strategies to balance knowledge retention and efficiency [54, 55]. Regularization-based methods [1, 15, 58] constrain updates to parameters deemed important for previously learned classes, thereby preserving historical knowledge without explicit data replay [52, 53, 60]. Isolation-based methods [26, 27, 33] allocate distinct model capacities for different tasks through pruning, masking, or gating, effectively mitigating interference among tasks. To further enhance representational flexibility, network expansion methods [40, 47, 65] dynamically increase the model’s capacity, expanding the feature extractor as new tasks arrive to accommodate emerging knowledge.

Despite their effectiveness, most of these approaches rely on extensive retraining or growing architectures, motivating the shift toward parameter-efficient adaptation that leverages frozen pre-trained models as stable knowledge bases for scalable CIL.

2.2. Pre-Trained Model-Based CIL

Leveraging the strong generalization ability of Pre-Trained Models (PTMs) [51], PTM-based Class-Incremental Learning (CIL) has recently emerged as a dominant paradigm [41, 46]. These methods shift the focus from training networks from scratch to adapting powerful PTMs through parameter-efficient fine-tuning (PEFT), aiming to retain generalization while continually learning new classes.

Early explorations adopted fine-tuning and projection-based strategies. SLCA [62] stabilizes representations by coupling a slowly updated backbone with a fast-learning classifier refined via feature generation. RanPAC [28] offers a training-free alternative by freezing the PTM and accumulating class prototypes through a random-projection layer, providing a strong rehearsal-free baseline. Prompt-based tuning methods then advanced PTM utilization by introducing learnable, task-adaptive prompts. L2P [44] employs an instance-conditioned prompt pool for dynamic retrieval,

while DualPrompt [43] separates prompts into general and expert components to capture both task-invariant and task-specific knowledge. CODA-Prompt [35] further enhances flexibility through modular prompt composition guided by attention mechanisms. Recent developments increasingly adopt adapter-based tuning to better balance adaptability and stability. APER+Adapter [66] combines frozen PTMs with prototype-based classification and branch merging to enhance generalization. LAE [9] introduces EMA-based calibration and integrates PEFT for adaptive prediction. EASE [67] learns lightweight adapters to span a scalable subspace and employs semantic prototype completion to mitigate cross-task interference. MOS [37] alleviates both parameter- and retrieval-level forgetting via adapter merging and self-refined retrieval, while SSIAT [38] maintains a shared adapter updated with semantic drift awareness and retrain the classifier through prototype sampling to reinforce stability-plasticity trade-offs.

Despite these advances, most PTM-based CIL methods focus on task-specific prompts or adapters and lack explicit mechanisms to guide representation structuring across tasks. In contrast, RSIAT bridges this gap by maintaining a single shared adapter and explicitly steering representations throughout incremental learning, achieving both compactness and cross-task coherence.

3. Preliminaries

3.1. Class-Incremental Learning

We introduce the definition of CIL. Data arrive in tasks $\{\mathcal{D}_t\}_{t=1}^T$, where $\mathcal{D}_t = \{(x_i^t, y_i^t)\}_{i=1}^{n_t}$ and $y_i^t \in \mathcal{Y}_t$. Label spaces are disjoint across tasks, $\mathcal{Y}_t \cap \mathcal{Y}_{t'} = \emptyset$ for $t \neq t'$. At task t , the model can only access $\mathcal{D}_t$ during training, and is evaluated on the cumulative label space $\mathcal{Y}_{1:t} = \mathcal{Y}_1 \cup \mathcal{Y}_2 \cup \dots \cup \mathcal{Y}_t$. We hope to find a model $\mathbf{f}(\mathbf{x}) : X \rightarrow \mathcal{Y}_b$ that minimizes the expected risk:

$$\mathbf{f}^* = \arg \min_{\mathbf{f} \in \mathcal{H}} \mathbb{E}_{(x,y) \sim \mathcal{D}_1 \cup \dots \cup \mathcal{D}_t} \mathbb{I}(y \neq \mathbf{f}(x)), \quad (1)$$

where $\mathcal{H}$ is the hypothesis space and $\mathbb{I}(\cdot)$ denotes the indicator function. Following typical PTM-based CIL works [67], we decouple a pre-trained backbone into a frozen feature extractor $\phi : \mathbb{R}^D \rightarrow \mathbb{R}^d$ and a classifier $W \in \mathbb{R}^{d \times |\mathcal{Y}_{1:t}|}$:

$$\mathbf{f}(x) = W^\top \phi(x), \quad W = [w_1, \dots, w_{|\mathcal{Y}_{1:t}|}]. \quad (2)$$

In PTM-based CIL, PEFT adapts a limited number of trainable modules (e.g., adapters, prompts) while keeping the pre-trained model frozen, which preserves prior knowledge and ensures parameter-efficient learning.

3.2. Baseline

Incremental adapter tuning. Our baseline builds directly upon SSIAT [38], which performs incremental adapter tun-

ing with a shared adapter A , while keeping the PTM backbone frozen. Given a batch $\mathcal{B}_t = \{(x_i, y_i)\}_t$ of samples from the new class set $\mathcal{Y}_t$, they extract embeddings $f_i^t = \phi(x_i; A^t) \in \mathbb{R}^d$ and the ℓ_2 -normalized features $\hat{z}_i^t = f_i^t / \|f_i^t\|$. Note that, during task t , only the adapter A^t and the classifier are trainable. Let the classifier weights be $W^t = [\mathbf{w}_c]_{c \in \mathcal{Y}_t}$, $\cos \theta_i^c = \hat{\mathbf{w}}_c^\top \hat{z}_i^t$ with $\hat{\mathbf{w}}_c = \mathbf{w}_c / \|\mathbf{w}_c\|$, a cosine loss is engaged [30],

$$\mathcal{L}_{\cos}^{(t)} = -\frac{1}{|\mathcal{B}_t|} \sum_{(x_i, y_i) \in \mathcal{B}_t} \log \frac{\exp(s(\cos \theta_i^{y_i} - m))}{\exp(s(\cos \theta_i^{y_i} - m)) + \sum_{c \in \mathcal{Y}_t \setminus \{y_i\}} \exp(s \cos \theta_i^c)}, \quad (3)$$

where s and m are the fixed scale and margin hyperparameters, respectively, enforcing tighter decision boundaries.

Semantic shift. After each task, we store the prototype and covariance of every new class as frozen statistics. This adds no trainable parameters and keeps the inference-time model size unchanged. However, as the shared adapter continues to be tuned on later tasks, features of earlier classes inevitably drift. Following prior work [38, 57], we estimate this drift by comparing embeddings of the current task’s data before and after adaptation, treating the shift as local variation to update older class prototypes. With updated prototypes and covariances, we model each class as a Gaussian-like distribution to synthesize features, which are then used to retrain the unified classification head. However, using only Eq. 3 during adapter tuning imposes weak geometric constraints, resulting in overlapping clusters, inspiring our representation steering.

4. Methodology

We propose a Representation-Steered Incremental Adapter Tuning (RSIAT) method, which is built upon the baseline. The framework of RSIAT is shown in Fig. 2. Its central idea is to balance stability and plasticity through a two-stage, continuous mechanism: (i) in the base task, we shape the representations to be compact within classes and well-separated between classes, allowing subsequent tasks to begin training from a more easily maintained representation structure; (ii) in the incremental task, **we use a residual autoencoder (RAE) to align the evolution space**, thus avoiding optimization resistance from strong regularization, while imposing moderate orthogonality constraints to prevent new features from overlapping with old prototypes. We consistently use a shared adapter to avoid parameter growth and module retrieval during inference.

4.1. Base Representation Steering

Representation steering. In the base task, the frozen PTM backbone network with shared adapter needs to be adapted




Figure 2. Overview of RSIAT. **Left:** Task 1 trains a shared adapter with a cosine classifier and a warm-up representation-steering loss $\mathcal{L}_{RS}$, which compacts within-class features and enlarges inter-class margins. **Right:** Task $t > 1$ updates the same adapter while a residual autoencoder projector P^t projects features from the previous model to the current space. The alignment loss $\mathcal{L}_{align}$ and the orthogonal loss $\mathcal{L}_{orth}$ align the feature space across tasks and steer the learning of new classes, while preserving the representations of old classes. The Projector is a lightweight autoencoder with an identity skip (residual).

to the dataset. We use the loss in Eq. 3 to train the adapter A^1 and classifier W^1 . As shown in Fig. 2, although this loss can help the model learn new classes, it cannot guide the model to construct a well-structured representation space. To actively construct a good feature space, we introduce a pairwise **representation-steering (RS) loss**. This loss maintains the transferability of PTM while promoting intra-class compactness and inter-class separability, enabling subsequent tasks to be trained on a feature space that is both robust and adaptive. RS loss can be decomposed into a positive loss ($\mathcal{L}_{pos}$) and a negative loss ($\mathcal{L}_{neg}$). The positive loss encourages same-class features to cluster, while the negative loss encourages different-class features to separate. Let $\hat{\mathbf{z}}_i^1$ denote the i -th normalized feature in the batch, $S_{ij} = \langle \hat{\mathbf{z}}_i^1, \hat{\mathbf{z}}_j^1 \rangle$ be the cosine similarity between samples i and j in batch $\mathcal{B}_1$, and $M_{ij}^+ = \mathbb{I}(y_i = y_j, i \neq j)$ and $M_{ij}^- = \mathbb{I}(y_i \neq y_j)$ be the positive and negative pair masks, respectively. The RS loss can be written as:

$$\begin{aligned}\mathcal{L}_{pos} &= \frac{1}{|P|} \sum_{i,j} (1 - S_{ij}) M_{ij}^+, \\ \mathcal{L}_{neg} &= \frac{1}{|N|} \sum_{i,j} (1 + S_{ij}) M_{ij}^-, \\ \mathcal{L}_{RS} &= \mathcal{L}_{pos} + \alpha \mathcal{L}_{neg},\end{aligned}\quad (4)$$

where $P = \{(i, j) \mid M_{ij}^+ = 1\}$ and $N = \{(i, j) \mid M_{ij}^- = 1\}$ collect the indices of positive and negative pairs, respectively; α balances the contributions of intra-class compactness and inter-class separability.

Warm-up schedule. Applying a strong RS loss at the very beginning undermines the effect of $\mathcal{L}_{cos}^{(1)}$ and hinders

the adapter’s adaptation to the dataset (as evidenced in Sec. 5.3). Therefore, we first warm up the RS coefficient $\lambda_{RS}(e)$ so the adapter can learn a coarse understanding of the dataset, and then gradually increase the strength of representation steering. Specifically, the RS coefficient is defined as:

$$\lambda_{RS}(e) = \lambda_{RS} \cdot \min\left(1, \frac{e}{E_w}\right), \quad (5)$$

where λ_{RS} is its target value after warm-up; e is the current epoch index; E_w is the warm-up duration (in epochs); and $\min(1, e/E_w)$ is the linear warm-up factor that grows from 0 to 1 and then remains capped at 1 (so that $\lambda_{RS}(e) = \lambda_{RS}$ once $e \geq E_w$). The base task training loss can be written as:

$$\mathcal{L}^{(1)} = \mathcal{L}_{cos}^{(1)} + \lambda_{RS}(e) \mathcal{L}_{RS}. \quad (6)$$

4.2. Incremental Representation Steering

RAE alignment. Updating the shared adapter for task t inevitably perturbs the representation learned at task $t - 1$. Inspired by optimal transport theory [20, 31], we align the $t - 1$ and t spaces by minimizing an alignment loss with a learnable projector. Prior PTM-based CIL studies report that strong regularizers (e.g., knowledge distillation) can slow or even impede the adapter’s learning [38].

To avoid this problem while enforcing alignment, we adopt a residual autoencoder (RAE) that learns a nonlinear projection while retaining an identity skip path, as shown by the projector on the right side of Fig. 2. We fix the down-sampling dimension to 64 while adjusting the middle up-sampling dimension to control the number of trainable parameters. Let $f^{t-1} = \phi(x; A^{t-1})$ be the features from the frozen $(t - 1)$ -th model and $f^t = \phi(x; A^t)$ be the features

Table 1. Experimental results on six CIL datasets with **ViT-B/16-IN21K** as the backbone. We report average accuracy $\bar{A}$ over tasks and final accuracy A_B . ‘IN-R/A’ stands for ‘ImageNet-R/A’, and ‘OmniBench’ stands for ‘OmniBenchmark’. All methods are implemented without using exemplars. **Best performance is shown in bold**, and second-best performance is shown in underline.

Method	CIFAR <i>B0I10</i>		IN-R <i>B0I20</i>		IN-A <i>B0I20</i>		OmniBench <i>B0I30</i>		CUB <i>B0I10</i>		VTAB <i>B0I10</i>	
	$\bar{A}$	A_B	$\bar{A}$	A_B	$\bar{A}$	A_B	$\bar{A}$	A_B	$\bar{A}$	A_B	$\bar{A}$	A_B
L2P [44]	89.23	84.29	75.46	69.77	49.39	41.71	73.36	64.69	67.05	56.25	77.11	77.10
DualPrompt [43]	87.36	82.30	73.10	67.18	53.71	41.67	73.92	65.52	77.47	66.54	83.36	81.23
CODA-Prompt [35]	91.30	86.85	77.97	72.27	53.54	42.73	77.03	68.09	84.00	73.37	83.90	83.02
RanPAC [28]	94.35	90.86	82.98	77.94	69.32	61.82	<u>85.95</u>	79.55	93.13	89.40	92.56	91.83
SLCA [62]	94.09	91.53	81.17	77.00	68.66	58.74	82.80	74.10	89.51	82.19	90.94	90.76
APER+Adapter [66]	90.94	85.75	75.82	67.95	60.47	49.37	80.75	74.37	92.21	86.73	85.95	84.35
EASE [67]	92.35	87.76	81.74	76.17	65.34	55.04	81.11	74.85	92.23	86.81	93.61	<u>93.55</u>
MOS [37]	<u>94.75</u>	<u>91.70</u>	82.96	77.93	69.13	59.12	85.91	<u>80.05</u>	<u>93.49</u>	<u>90.12</u>	92.62	<u>92.79</u>
SSIAT [38]	94.35	91.35	<u>83.63</u>	<u>79.38</u>	<u>70.83</u>	<u>62.43</u>	84.31	<u>77.63</u>	93.38	89.53	<u>94.21</u>	91.28
RSIAT (Ours)	95.15	92.20	86.92	82.75	74.89	66.23	86.42	81.27	93.99	90.59	94.65	94.69

from the t -th model. The projector is defined as:

$$P^t(f^{t-1}) = \text{AutoEncoder}^t(f^{t-1}) + f^{t-1}. \quad (7)$$

Because the adapter and the projector are co-trained, the projector must continually track the model’s evolving representation. At the start of task t , we initialize $A^t = A^{t-1}$ and the projector to zero. Under the residual form, the projector is effectively an identity mapping (i.e., an I -projection), so the projection is perfectly aligned with the model at initialization. As optimization proceeds and the adapter updates, the projector only needs to model the incremental change of the representation on top of the identity path, it learns to account for the drift from f^{t-1} to f^t without imposing constraints on the adapter. This design lets the projector keep up with representation shifts while avoiding any hand-brake effect on the adapter’s parameter updates. As shown in Fig. 2, the projector projects f^{t-1} and aligns them with f^t . The L_2 alignment loss used can be expressed as:

$$\mathcal{L}_{\text{align}}^{(t)} = \frac{1}{|\mathcal{B}_t|} \sum_{x \in \mathcal{B}_t} \|f^t - P^t(f^{t-1})\|_2^2. \quad (8)$$

We use the $\mathcal{L}_{\text{align}}^{(t)}$ and RAE projector to align the gradually evolving feature space and avoid the optimization resistance caused by strong regularization.

Orthogonal representation steering. To reduce representational conflicts between the old and new classes, we use the memory of old class prototypes $\{\mu_c\}_{c \in \mathcal{Y}_{1:t-1}}$ accumulated at previous tasks. We map both the old prototypes and the old-tasks features into the current space with the projector $P^t(\cdot)$ and ℓ_2 -normalize them, $\hat{p}_c = P^t(\mu_c)/\|P^t(\mu_c)\|$ and $\hat{u}_i = P^t(f_i^{t-1})/\|P^t(f_i^{t-1})\|$. We then calculate the cosine similarity among these normalized vectors to encourage them to be orthogonal, thus ensuring that learning new classes does not affect the representation of old classes. The

orthogonal loss is expressed as follows:

$$\mathcal{L}_{\text{orth}}^{(t)} = \frac{1}{|\mathcal{Y}_{1:t-1}| |\mathcal{B}_t|} \sum_{c=1}^{|\mathcal{Y}_{1:t-1}|} \sum_{i=1}^{|\mathcal{B}_t|} |\langle \hat{p}_c, \hat{u}_i \rangle|, \quad (9)$$

where $|\langle \hat{p}_c, \hat{u}_i \rangle|$ represents the absolute value of the cosine similarity between the prototype and the feature. $\mathcal{L}_{\text{orth}}^{(t)}$ steers the separation of new classes from the prototypes of old classes, learning new classes with minimal impact on old classes, thus enhancing the model’s stability.

The incremental task training loss can be written as:

$$\mathcal{L}^{(t)} = \mathcal{L}_{\text{cos}}^{(t)} + \beta \mathcal{L}_{\text{align}}^{(t)} + \gamma \mathcal{L}_{\text{orth}}^{(t)}. \quad (10)$$

Inference procedure. During inference, the residual autoencoder is discarded. RSIAT infers with the frozen PTM backbone, the shared adapter, and the classifier, keeping the inference-time parameter count strictly constant. We summarize the pseudo code of RSIAT in Appendix A.

5. Experiments

We evaluate RSIAT on six CIL datasets against state-of-the-art methods under various CIL settings, and further verify its effectiveness with ablation studies, hyperparameter analyses, and t-SNE visualizations.

5.1. Implementation Details

Benchmark datasets. We evaluate the methods on six CIL benchmarks datasets, CIFAR100 [19], ImageNet-R [12], ImageNet-A [13], OmniBenchmark [63], CUB200 [39] and VTAB [59]. CIFAR100 contains 100 classes, ImageNet-R, ImageNet-A, and CUB200 contain 200 classes, OmniBenchmark contains 300 classes, and VTAB contains 50 classes. These six datasets have a large domain gap compared to the ImageNet dataset used for PTMs, making them suitable for use as evaluation datasets. Following the benchmark setting [67], we denote **BmIn** as m base classes and




Figure 3. Performance curve on five CIL datasets with *ViT-B/16-IN1K* as the backbone. All methods are implemented without using exemplars. *BmIn* denotes *m* base classes and *n* new classes per incremental task. Annotate the final incremental task with numbers to indicate RSIAT’s relative improvement over the runner-up method.

n newly added classes per incremental tasks. Class orders are randomly shuffled with seed 1993, and all methods share the same train/test partitions. More details about the datasets can be found in the Appendix B.

Baseline methods. We engage both PTM-Based CIL methods and traditional CIL methods as our baselines. PTM-based CIL methods include L2P [44], DualPrompt [43], CODA-Prompt [35], RanPAC [28], SLCA [62], APER+Adapter [66], MOS [37], EASE [67], and SSIAT [38]. Traditional CIL methods include iCaRL [32], DER [47], FOSTER [40], MEMO [65], and TagFex [64]. Please refer to Appendix C for more details.

Training configurations. We run all experiments on NVIDIA 4090 and reproduce the compared methods with PyTorch [29] within the PILOT framework [36]. We use ViT-B/16 [6] as the PTM and consider two backbones: *ViT-B/16-IN21K* and *ViT-B/16-IN1K*. Both are initialized from ImageNet-21K pretraining, while the latter is further fine-tuned on ImageNet-1K. Following the work [38], we set the adapter dimension as 64 and use the cosineAnneal scheduler. We set different parameters for different datasets. The specific parameters are detailed in the Appendix D.

Evaluation metrics. We report two standard metrics: final accuracy [24] and average accuracy [32]. Given accuracies $a_{k,j} \in [0, 1]$ measured on the j -th task ($j \leq k$) after training tasks from 1 to k . The final accuracy, $A_B = \frac{1}{B} \sum_{j=1}^B a_{B,j}$, reflects the performance after learning B tasks. The average

Table 2. Comparison to traditional CIL methods with *ViT-B/16-IN21K* as the backbone. Baselines use 20 exemplars per class.

Method	Exemplars	IN-R B0I20		CIFAR B0I10	
		$\bar{A}$	A_B	$\bar{A}$	A_B
iCaRL [32]	20 / class	72.42	60.67	82.46	73.87
DER [47]	20 / class	80.48	74.32	86.04	77.93
FOSTER [40]	20 / class	81.34	74.48	89.87	84.91
MEMO [65]	20 / class	74.80	66.62	84.08	75.79
TagFex [64]	20 / class	83.23	78.45	92.17	89.26
RSIAT	0	86.92	82.75	95.15	92.20

accuracy, $\bar{A} = \frac{1}{B} \sum_{k=1}^B A_k = \frac{1}{B} \sum_{k=1}^B (\frac{1}{k} \sum_{j=1}^k a_{k,j})$, summarizes performance across all stages and captures how well the model maintains accuracy as classes accumulate. Higher A_B and $\bar{A}$ indicate superior continual learning.

5.2. Overall Evaluation

We compare RSIAT with PTM-based and traditional CIL methods using two backbone initializations across multiple incremental settings.

Tab. 1 reports results on six CIL datasets with *ViT-B/16-IN21K* as the pre-trained backbone [37]. RSIAT achieves the best average accuracy $\bar{A}$ and the best final accuracy A_B on all six datasets. The advantages are pronounced on ImageNet-A/R, indicating that RSIAT performs well across diverse data domains and exhibits strong robustness. Fig. 3 shows accuracy curves under the *ViT-B/16-IN1K* pre-trained backbone. RSIAT is competitive at early stages and pro-




Figure 4. Results of IN-R/A in CIL settings with large base classes and long sequence tasks. All methods use the same backbone *ViT-B/16-IN21K*. Annotate the final incremental task with numbers to indicate RSIAT’s relative improvement over the runner-up method.

Table 3. Ablation study on ImageNet-R and ImageNet-A datasets.

Method	IN-R <i>B0I20</i>		IN-A <i>B0I20</i>	
	$\bar{A}$	A_B	$\bar{A}$	A_B
Baseline	83.63	79.38	70.83	62.43
+ $\mathcal{L}_{orth}$	84.15	78.32	72.83	61.95
+ $\mathcal{L}_{align}$	84.72	80.17	74.30	65.04
+ IRS	84.86	80.52	74.37	65.31
+ IRS (<i>wo_res</i>)	80.59	74.02	65.37	51.48
+ IRS + $\mathcal{L}_{RS}$	85.36	81.45	74.62	65.50
RSIAT	86.92	82.75	74.89	66.23

gressively widens its lead as classes accumulate, reflecting our designs: (i) shaping compact within-class structure and separated inter-class geometry at the base stage, and (ii) aligning spaces and encouraging separation between new and old features in later stages. The final-stage annotations mark our improvement over the runner-up method: +1.91 on ImageNet-A and +1.47 on ImageNet-R; other datasets also show improvements to varying degrees.

Following the setting in [42], we compare RSIAT with the traditional CIL methods combined with the *ViT-B/16-IN21K* backbone. It is worth noting that RSIAT is exemplar-free, whereas the engaged baselines are exemplar-based methods. As shown in Tab. 2, RSIAT is competitive even when the baselines engage 20 examples pre-class. Besides, we conduct large base classes and long sequence tasks on ImageNet-R/A by using *ViT-B/16-IN21K* as the pre-trained backbone. As shown in Fig. 4, RSIAT (i) starts competitively at early increments, (ii) degrades more gracefully as the number of classes grows, and (iii) achieves the highest last-stage accuracy. It confirms that RSIAT adapts well to diverse experimental settings. More experimental results are provided in the Appendix E.1 & E.2.

5.3. Ablation Study

We ablate each RSIAT component on ImageNet-R and ImageNet-A in Tab. 3. Starting from the Baseline (cosine head and shared adapter only), independently adding the orthogonal loss (+ $\mathcal{L}_{orth}$) improves average accuracy ($\bar{A}$) by preventing new features from overlapping with old prototypes, though it slightly restricts final plasticity (A_B). Al-




Figure 5. Hyperparameter study on ImageNet-R (*B0I20*). Heatmap of average accuracy $\bar{A}$ versus projector dimension d_p (x-axis) and adapter bottleneck d_a (y-axis).

ternatively, introducing the alignment loss (+ $\mathcal{L}_{align}$) aligns the current embedding space with the projected latent space to explicitly mitigate feature drift, yielding consistent gains across both metrics. Combining them into Incremental Representation Steering (+ IRS) provides the optimal synergy, effectively keeping old and new concepts separated. The + IRS (*wo_res*) variant removes the projector’s residual identity path, causing a severe performance drop (e.g., A_B falls to 51.48% on IN-A). This confirms that without the identity prior, the projector over-regularizes the adapter and hinders the acquisition of new knowledge. Integrating Base Representation Steering (+ IRS + $\mathcal{L}_{RS}$) brings further gains, proving that a compact base feature space seeds stable downstream adaptation. Finally, the full RSIAT introduces a warm-up schedule. The notable improvement over the non-warmed-up variant validates the need to initially carve margins before tightening clusters, ultimately achieving the most favorable stability-plasticity trade-off.

5.4. Further Analysis

Hyperparameter study. RSIAT contains two hyperparameters that affect the number of trainable parameters: the adapter bottleneck dimension d_a and the residual projector code dimension d_p . We sweep $d_a \in \{16, 32, 64, 128, 256\}$ and $d_p \in \{128, 256, 384, 512, 768\}$ on ImageNet-R (*B0I20*) and report the average accuracy $\bar{A}$ in Fig. 5. The results show that the performance surface is relatively flat around a ridge near $d_p \approx$




Figure 6. t-SNE embeddings across tasks and methods. Columns correspond to Task 1/2/3, rows correspond to MOS, RanPAC, SSIAT, and RSIAT (ours). Numbers in each panel denote Recall@1 measured in the original feature space before t-SNE. RSIAT forms compact clusters and preserves separation across tasks.

384 and $d_a \in [64, 256]$. Although the absolute peak occurs at $(d_p=384, d_a=256)$ with $\bar{A} \approx 87.04$, using $(d_p=384, d_a=64)$ yields $\bar{A} \approx 86.92$ (only ~ 0.12 lower) while reducing the adapter parameter count to roughly 1/4. Given the small performance gap, we adopt $(d_p=384, d_a=64)$ as the default to minimize extra parameters with minimal accuracy sacrifice.

Visualizations. We analyze how different methods shape the feature space across three incremental tasks (Task 1/2/3) under the *ViT-B/16-IN21K* backbone. At each task, we visualize the features with t-SNE [25] and report the nearest-neighbor class-match rate Recall@1 [48] in the original feature space to quantify within-class compactness and between-class dispersion. As shown in Fig. 6, RSIAT produces the most compact and well-separated clusters at all three tasks (Task 1/2/3) and achieves the highest and most stable Recall@1 scores ($94.5\% \rightarrow 87.4\% \rightarrow 84.3\%$), with a markedly smaller drop than MOS, RanPAC, and SSIAT. It indicates that RSIAT learns and maintains strong represen-

tations (tighter within-class cohesion and larger between-class margins) while exhibiting smaller cross-task prototype drift. More t-SNE results are given in the Appendix E.3.

6. Conclusion

We introduced Representation-Steered Incremental Adapter Tuning (RSIAT), a unified framework for parameter-efficient class-incremental learning. By steering representation learning in the base task and employing a residual autoencoder to align features across tasks, RSIAT achieves consistent and discriminative representations without increasing model parameters. This design effectively balances stability and plasticity, eliminating the linear parameter growth common in prior PTM-based approaches. Extensive experiments across diverse benchmarks validate RSIAT’s efficiency and robustness, offering a scalable solution for continual adaptation in large pre-trained models. Future work will explore broader modalities and further structural guidance for lifelong representation learning.

Acknowledgment

This work was supported by the National Natural Science Foundation of China (Grants 62476264 and 62406312).

References

- [1] Rahaf Aljundi, Francesca Babiloni, Mohamed Elhoseiny, Marcus Rohrbach, and Tinne Tuytelaars. Memory aware synapses: Learning what (not) to forget. In *Proceedings of the European conference on computer vision (ECCV)*, pages 139–154, 2018. 1, 2
- [2] Zhulin An, Xinqiang Yu, Chu Wang, Yinlong Zhang, and Chunhe Song. Embodied intelligence: Recent advances and future perspectives. *The Innovation Informatics*, 1(1): 100008–1, 2025. 1
- [3] Jihwan Bang, Heesu Kim, YoungJoon Yoo, Jung-Woo Ha, and Jonghyun Choi. Rainbow memory: Continual learning with a memory of diverse samples. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, pages 8218–8227, 2021. 1, 2
- [4] Arslan Chaudhry, Puneet K Dokania, Thalaiyasingam Ajanthan, and Philip HS Torr. Riemannian walk for incremental learning: Understanding forgetting and intransigence. In *Proceedings of the European conference on computer vision (ECCV)*, pages 532–547, 2018. 1, 2
- [5] Shoufa Chen, Chongjian Ge, Zhan Tong, Jiangliu Wang, Yibing Song, Jue Wang, and Ping Luo. Adaptformer: Adapting vision transformers for scalable visual recognition. *Advances in Neural Information Processing Systems*, 35:16664–16678, 2022. 1
- [6] Alexey Dosovitskiy. An image is worth 16x16 words: Transformers for image recognition at scale. *arXiv preprint arXiv:2010.11929*, 2020. 6
- [7] Arthur Douillard, Matthieu Cord, Charles Ollion, Thomas Robert, and Eduardo Valle. Podnet: Pooled outputs distillation for small-tasks incremental learning. In *European Conference on Computer Vision*, pages 86–102. Springer, 2020. 2
- [8] Robert M French. Catastrophic forgetting in connectionist networks. *Trends in cognitive sciences*, 3(4):128–135, 1999. 1
- [9] Qiankun Gao, Chen Zhao, Yifan Sun, Teng Xi, Gang Zhang, Bernard Ghanem, and Jian Zhang. A unified continual learning framework with general parameter-efficient tuning. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pages 11483–11493, 2023. 3
- [10] Zhanxin Gao, Jun Cen, and Xiaobin Chang. Consistent prompting for rehearsal-free continual learning. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 28463–28473, 2024. 2
- [11] Xu Han, Zhengyan Zhang, Ning Ding, Yuxian Gu, Xiao Liu, Yuqi Huo, Jiezhong Qiu, Yuan Yao, Ao Zhang, Liang Zhang, et al. Pre-trained models: Past, present and future. *AI Open*, 2:225–250, 2021. 1
- [12] Dan Hendrycks, Steven Basart, Norman Mu, Saurav Kada-vath, Frank Wang, Evan Dorundo, Rahul Desai, Tyler Zhu, Samyak Parajuli, Mike Guo, et al. The many faces of robustness: A critical analysis of out-of-distribution generalization. In *Proceedings of the IEEE/CVF international conference on computer vision*, pages 8340–8349, 2021. 5, 1
- [13] Dan Hendrycks, Kevin Zhao, Steven Basart, Jacob Steinhardt, and Dawn Song. Natural adversarial examples. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, pages 15262–15271, 2021. 5, 1
- [14] Libo Huang, Zhulin An, Yan Zeng, Yongjun Xu, et al. Kfc: Knowledge reconstruction and feedback consolidation enable efficient and effective continual generative learning. In *The Second Tiny Papers Track at ICLR 2024*, 2024. 1
- [15] Libo Huang, Yan Zeng, Chuanguang Yang, Zhulin An, Boyu Diao, and Yongjun Xu. etag: Class-incremental learning via embedding distillation and task-oriented generation. In *Proceedings of the AAAI Conference on Artificial Intelligence*, pages 12591–12599, 2024. 1, 2
- [16] Libo Huang, Zhulin An, Chuanguang Yang, Boyu Diao, Fei Wang, Yan Zeng, Zhifeng Hao, and Yongjun Xu. Preprompt: Predictive prompting for class incremental learning. *arXiv preprint arXiv:2505.08586*, 2025. 1
- [17] Menglin Jia, Luming Tang, Bor-Chun Chen, Claire Cardie, Serge Belongie, Bharath Hariharan, and Ser-Nam Lim. Visual prompt tuning. In *European conference on computer vision*, pages 709–727. Springer, 2022. 1
- [18] James Kirkpatrick, Razvan Pascanu, Neil Rabinowitz, Joel Veness, Guillaume Desjardins, Andrei A Rusu, Kieran Milan, John Quan, Tiago Ramalho, Agnieszka Grabska-Barwinska, et al. Overcoming catastrophic forgetting in neural networks. *Proceedings of the national academy of sciences*, 114(13):3521–3526, 2017. 1, 5
- [19] Alex Krizhevsky, Geoffrey Hinton, et al. Learning multiple layers of features from tiny images. 2009. 5, 1
- [20] Qiwei Li, Yuxin Peng, and Jiahuan Zhou. Fcs: Feature calibration and separation for non-exemplar class incremental learning. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, pages 28495–28504, 2024. 4
- [21] Xiangqi Li, Libo Huang, Zhulin An, Weilun Feng, Chuanguang Yang, Boyu Diao, Fei Wang, and Yongjun Xu. Geometric feature embedding for effective 3d few-shot class incremental learning. In *Forty-second International Conference on Machine Learning*, 2025. 1
- [22] Xiangqi Li, Libo Huang, Jiarui Zhao, Weilun Feng, Chuanguang Yang, Zhulin An, and Yongjun Xu. Point-UQ: An uncertainty-quantification paradigm for point cloud few-shot class incremental learning. In *The Fourteenth International Conference on Learning Representations*, 2026. 1
- [23] Zhizhong Li and Derek Hoiem. Learning without forgetting. *IEEE transactions on pattern analysis and machine intelligence*, 40(12):2935–2947, 2017. 1
- [24] David Lopez-Paz and Marc’Aurelio Ranzato. Gradient episodic memory for continual learning. *Advances in neural information processing systems*, 30, 2017. 1, 6
- [25] Laurens van der Maaten and Geoffrey Hinton. Visualizing data using t-sne. *Journal of machine learning research*, 9 (Nov):2579–2605, 2008. 8

- [26] Arun Mallya and Svetlana Lazebnik. Packnet: Adding multiple tasks to a single network by iterative pruning. In *Proceedings of the IEEE conference on Computer Vision and Pattern Recognition*, pages 7765–7773, 2018. 1, 2
- [27] Arun Mallya, Dillon Davis, and Svetlana Lazebnik. Piggyback: Adapting a single network to multiple tasks by learning to mask weights. In *Proceedings of the European conference on computer vision (ECCV)*, pages 67–82, 2018. 1, 2
- [28] Mark D McDonnell, Dong Gong, Amin Parvaneh, Ehsan Abbasnejad, and Anton Van den Hengel. Ranpac: Random projections and pre-trained models for continual learning. *Advances in Neural Information Processing Systems*, 36:12022–12053, 2023. 2, 5, 6, 4
- [29] Adam Paszke, Sam Gross, Francisco Massa, Adam Lerer, James Bradbury, Gregory Chanan, Trevor Killeen, Zeming Lin, Natalia Gimelshein, Luca Antiga, et al. Pytorch: An imperative style, high-performance deep learning library. *Advances in neural information processing systems*, 32, 2019. 6
- [30] Can Peng, Kun Zhao, Tianren Wang, Meng Li, and Brian C Lovell. Few-shot class-incremental learning from an open-set perspective. In *European Conference on Computer Vision*, pages 382–397. Springer, 2022. 3
- [31] Gabriel Peyré, Marco Cuturi, et al. Computational optimal transport: With applications to data science. *Foundations and Trends® in Machine Learning*, 11(5-6):355–607, 2019. 4
- [32] Sylvestre-Alvise Rebuffi, Alexander Kolesnikov, Georg Sperl, and Christoph H Lampert. icarl: Incremental classifier and representation learning. In *Proceedings of the IEEE conference on Computer Vision and Pattern Recognition*, pages 2001–2010, 2017. 1, 2, 6
- [33] Joan Serra, Didac Suris, Marius Miron, and Alexandros Karatzoglou. Overcoming catastrophic forgetting with hard attention to the task. In *International conference on machine learning*, pages 4548–4557. PMLR, 2018. 1, 2
- [34] Wuxuan Shi and Mang Ye. Prospective representation learning for non-exemplar class-incremental learning. *Advances in Neural Information Processing Systems*, 37:995–1018, 2024. 1
- [35] James Seale Smith, Leonid Karlinsky, Vyshnavi Gutta, Paola Cascante-Bonilla, Donghyun Kim, Assaf Arbelle, Rameswar Panda, Rogerio Feris, and Zsolt Kira. Coda-prompt: Continual decomposed attention-based prompting for rehearsal-free continual learning. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, pages 11909–11919, 2023. 3, 5, 6, 2, 4
- [36] Hai-Long Sun, Da-Wei Zhou, De-Chuan Zhan, and Han-Jia Ye. Pilot: A pre-trained model-based continual learning toolbox, 2025. 6
- [37] Hai-Long Sun, Da-Wei Zhou, Hanbin Zhao, Le Gan, De-Chuan Zhan, and Han-Jia Ye. Mos: Model surgery for pre-trained model-based class-incremental learning. In *Proceedings of the AAAI Conference on Artificial Intelligence*, pages 20699–20707, 2025. 3, 5, 6, 2, 4
- [38] Yuwen Tan, Qinzhao Zhou, Xiang Xiang, Ke Wang, Yuchuan Wu, and Yongbin Li. Semantically-shifted incremental adapter-tuning is a continual vitransformer. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 23252–23262, 2024. 3, 4, 5, 6, 2, 8
- [39] Catherine Wah, Steve Branson, Peter Welinder, Pietro Perona, and Serge Belongie. The caltech-ucsd birds-200-2011 dataset. 2011. 5, 1
- [40] Fu-Yun Wang, Da-Wei Zhou, Han-Jia Ye, and De-Chuan Zhan. Foster: Feature boosting and compression for class-incremental learning. In *European conference on computer vision*, pages 398–414. Springer, 2022. 2, 6
- [41] Yabin Wang, Zhiwu Huang, and Xiaopeng Hong. S-prompts learning with pre-trained transformers: An occam’s razor for domain incremental learning. *Advances in Neural Information Processing Systems*, 35:5682–5695, 2022. 2
- [42] Yan Wang, Da-Wei Zhou, and Han-Jia Ye. Integrating task-specific and universal adapters for pre-trained model-based class-incremental learning. *arXiv preprint arXiv:2508.08165*, 2025. 1, 7
- [43] Zifeng Wang, Zizhao Zhang, Sayna Ebrahimi, Ruoxi Sun, Han Zhang, Chen-Yu Lee, Xiaoqi Ren, Guolong Su, Vincent Perot, Jennifer Dy, et al. Dualprompt: Complementary prompting for rehearsal-free continual learning. In *European conference on computer vision*, pages 631–648. Springer, 2022. 3, 5, 6, 2, 4
- [44] Zifeng Wang, Zizhao Zhang, Chen-Yu Lee, Han Zhang, Ruoxi Sun, Xiaoqi Ren, Guolong Su, Vincent Perot, Jennifer Dy, and Tomas Pfister. Learning to prompt for continual learning. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, pages 139–149, 2022. 1, 2, 5, 6, 4
- [45] Fangwen Wu, Lechao Cheng, Shengeng Tang, Xiaofeng Zhu, Chaowei Fang, Dingwen Zhang, and Meng Wang. Navigating semantic drift in task-agnostic class-incremental learning. *arXiv preprint arXiv:2502.07560*, 2025. 2
- [46] Tz-Ying Wu, Gurumurthy Swaminathan, Zhizhong Li, Avinash Ravichandran, Nuno Vasconcelos, Rahul Bhotika, and Stefano Soatto. Class-incremental learning with strong pre-trained models. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 9601–9610, 2022. 2
- [47] Shipeng Yan, Jiangwei Xie, and Xuming He. Der: Dynamically expandable representation for class incremental learning. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, pages 3014–3023, 2021. 2, 6
- [48] Chuanguang Yang, Zhulin An, Linhang Cai, and Yongjun Xu. Knowledge distillation using hierarchical self-supervision augmented distribution. *IEEE transactions on neural networks and learning systems*, 35(2):2094–2108, 2022. 8
- [49] Chuanguang Yang, Helong Zhou, Zhulin An, Xue Jiang, Yongjun Xu, and Qian Zhang. Cross-image relational knowledge distillation for semantic segmentation. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, pages 12319–12328, 2022. 2
- [50] Chuanguang Yang, Zhulin An, Helong Zhou, Fuzhen Zhuang, Yongjun Xu, and Qian Zhang. Online knowledge

- distillation via mutual contrastive learning for visual recognition. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 45(8):10212–10227, 2023. 2
- [51] Chuanguang Yang, Zhulin An, Libo Huang, Junyu Bi, Xinqiang Yu, Han Yang, Boyu Diao, and Yongjun Xu. Clip-kd: An empirical study of clip model distillation. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 15952–15962, 2024. 2
- [52] Dongbao Yang, Yu Zhou, Wei Shi, Dayan Wu, and Weiping Wang. Rd-iod: Two-level residual-distillation-based triple-network for incremental object detection. *ACM Transactions on Multimedia Computing, Communications, and Applications (TOMM)*, 18(1):1–23, 2022. 2
- [53] Dongbao Yang, Yu Zhou, Aoting Zhang, Xurui Sun, Dayan Wu, Weiping Wang, and Qixiang Ye. Multi-view correlation distillation for incremental object detection. *Pattern Recognition*, 131:108863, 2022. 2
- [54] Dongbao Yang, Yu Zhou, Xiaopeng Hong, Aoting Zhang, and Weiping Wang. One-shot replay: Boosting incremental object detection via retrospecting one object. In *Proceedings of the AAAI Conference on Artificial Intelligence*, pages 3127–3135, 2023. 2
- [55] Dongbao Yang, Yu Zhou, Xiaopeng Hong, Aoting Zhang, Xin Wei, Linchengxi Zeng, Zhi Qiao, and Weiping Wang. Pseudo object replay and mining for incremental object detection. In *Proceedings of the 31st ACM International Conference on Multimedia*, pages 153–162, 2023. 2
- [56] Lu Yu, Bartłomiej Twardowski, Xialei Liu, Luis Herranz, Kai Wang, Yongmei Cheng, Shangling Jui, and Joost van de Weijer. Semantic drift compensation for class-incremental learning. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, pages 6982–6991, 2020. 2
- [57] Lu Yu, Bartłomiej Twardowski, Xialei Liu, Luis Herranz, Kai Wang, Yongmei Cheng, Shangling Jui, and Joost van de Weijer. Semantic drift compensation for class-incremental learning. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, pages 6982–6991, 2020. 3
- [58] Friedemann Zenke, Ben Poole, and Surya Ganguli. Continual learning through synaptic intelligence. In *International conference on machine learning*, pages 3987–3995. PMLR, 2017. 1, 2
- [59] Xiaohua Zhai, Joan Puigcerver, Alexander Kolesnikov, Pierre Ruyssen, Carlos Riquelme, Mario Lucic, Josip Djolonga, Andre Susano Pinto, Maxim Neumann, Alexey Dosovitskiy, et al. A large-scale study of representation learning with the visual task adaptation benchmark. *arXiv preprint arXiv:1910.04867*, 2019. 5, 2
- [60] Aoting Zhang, Dongbao Yang, Chang Liu, Xiaopeng Hong, Miao Shang, and Yu Zhou. Dca: Dividing and conquering amnesia in incremental object detection. In *Proceedings of the AAAI Conference on Artificial Intelligence*, pages 9851–9859, 2025. 2
- [61] Aoting Zhang, Dongbao Yang, Chang Liu, Xiaopeng Hong, and Yu Zhou. Specifying what you know or not for multi-label class-incremental learning. In *Proceedings of the AAAI Conference on Artificial Intelligence*, pages 22345–22353, 2025. 2
- [62] Gengwei Zhang, Liyuan Wang, Guoliang Kang, Ling Chen, and Yunchao Wei. Slca: Slow learner with classifier alignment for continual learning on a pre-trained model. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pages 19148–19158, 2023. 2, 5, 6, 4
- [63] Yuanhan Zhang, Zhenfei Yin, Jing Shao, and Ziwei Liu. Benchmarking omni-vision representation through the lens of visual realms. In *European conference on computer vision*, pages 594–611. Springer, 2022. 5, 1
- [64] Bowen Zheng, Da-Wei Zhou, Han-Jia Ye, and De-Chuan Zhan. Task-agnostic guided feature expansion for class-incremental learning. In *Proceedings of the Computer Vision and Pattern Recognition Conference*, pages 10099–10109, 2025. 1, 6, 2
- [65] Da-Wei Zhou, Qi-Wei Wang, Han-Jia Ye, and De-Chuan Zhan. A model or 603 exemplars: Towards memory-efficient class-incremental learning. *arXiv preprint arXiv:2205.13218*, 2022. 2, 6
- [66] Da-Wei Zhou, Zi-Wen Cai, Han-Jia Ye, De-Chuan Zhan, and Ziwei Liu. Revisiting class-incremental learning with pre-trained models: generalizability and adaptivity are all you need. *International Journal of Computer Vision*, 2024. 3, 5, 6, 1, 2, 4
- [67] Da-Wei Zhou, Hai-Long Sun, Han-Jia Ye, and De-Chuan Zhan. Expandable subspace ensemble for pre-trained model-based class-incremental learning. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 23554–23564, 2024. 3, 5, 6, 2, 4